

The logo consists of the word "ARUN'S" in a stylized font inside an oval border, with the word "INDEX" in large, bold, sans-serif letters below it.

Name Shashank. M.S. Subject Machine Learning

Standard 3rd Year BE Section CSE 6D section Roll No. IBM21CS301

School/ College B.M.S. College of Engineering

S. No.	Date	Title	Page No.	Teacher's Sign
17	21/03/24	Import & Export csv file LAB 1	1-2	10 21/03/24
27	28/03/24	LAB 2	3-7	10 28/03/24
37	04/04/24	LAB 3	8-9	10 4/4/24
47	18/04/24	LAB 4	10-12	10 18/04/24
57	25/04/24	LAB 5	13-14	10 25/04/24
67	09/05/24	LAB 6	15-17	10 9/5
77	23/05/24	LAB 7	18-21	10 23/05/24
87	30/05/24	LAB 8	22-24	10 30/05/24

Machine Learning

Lab - 1

- 1) Write a python program to import and export data using pandas library function.

Reading data from URL

```
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
```

```
col_names = ["Sepal-length-in-cm",  
             "Sepal-width-in-cm",  
             "petal-length-in-cm",  
             "petal-width-in-cm",  
             "class"]
```

```
iris_data = pd.read_csv(url, names=col_names)
```

```
iris_data.head()
```

Output:

	Sepal-length-in-cm	Sepal-width-in-cm	petal-length-in-cm	petal-width-in-cm	class
0	5.1	3.5	1.4	0.2	Iris-setosa

importing & displaying data

import pandas as pd

iris_data = pd.read_csv("C:/Users/Admin/Downloads/iris dataset")

iris_data.head

5.1 3.5 1.4 0.2 Iris setosa

0 4.9 3.0 1.4 0.2 Iris setosa

1 4.7 3.2 1.3 0.2 Iris setosa

Exporting file to current working directory

iris_data.to_csv("cleaned_iris_data.csv")

Voice
21/5/19

Machine Learning

Lab - 2

2) Get the data

```
import os  
import tarfile  
import urllib
```

DOWNLOAD_ROOT = "https://raw.githubusercontent.com/agronn/handson-ml2/master/"

HOUSING_PATH = os.path.join("data", "01")

HOUSING_URL = DOWNLOAD_ROOT + "datasets/housing/housing.tgz"

fetch_housing_data()

```
import pandas as pd  
def load_housing_data(housing_path=HOUSING_PATH):
```

data_path = os.path.join(housing_path, "housing.csv")
return pd.read_csv(data_path)

housing = load_housing_data()

housing.head()

housing.info()

3)

Discover and Visualize the Data to gain Insights

Strat-train-set.shape, Strat-test-set.shape

Strat-test-set.reset_index().to_feather(fname='data/01/strat-test')

housing = Strat-train-set.copy(); housing.shape

Now for visualizing geographical Data

housing.plot(kind='scatter', x='longitude', y='latitude')
plt.show()

housing.plot(kind='scatter', x='longitude', y='latitude', alpha=0.1)
plt.show()

Now for correlation

corr_matrix = housing.corr()

corr_matrix['median_house_value']

Sort Values (ascending=False)

from pandas.plotting import scatter_matrix

attributes = ['median_house_value', 'median_income']

scatter_matrix = housing[attributes], figsize=(12,8)
plt.show()

$$\text{choosing } [{}^e \text{ rooms_per_household?}] = \frac{\text{housing } [{}^e \text{ total rooms}]}{\text{housing } [{}^e \text{ households}]}$$

4) Prepare Data for Machine learning Algorithms

from sklearn.impute import SimpleImputer

imputer = SimpleImputer (strategy = 'median')

housing_num = housing.drop ('ocean_proximity', axis=1)

imputer.fit (housing_num)

imputer.statistics - housing_num.median ().values

from sklearn.preprocessing import OrdinalEncoder

ordinal_encoder = OrdinalEncoder ()

ordinal_encoder.categories

attr_adder = CombinedAttributeAdder (add_bedroom, per_room=False)

from sklearn.compose import column_transformer

num_attributes = housing_num.columns.to_list ()

5) Select and Train a Model

- from `sklearn.linear_model import LinearRegression`

`lin-reg = LinearRegression()`

`lin-reg.fit(x = housing['prepared'],
y = housing['labels'])`

`housing['predictions'] = lin-reg.predict(housing['prepared'])`

`lin-mse = mean_squared_error(housing['label'],
housing['predictions'])`

`lin-rmse = np.sqrt(lin-mse)`

`lin-rmse`

`tree-reg = DecisionTreeRegressor()`

`tree-reg.fit(x = housing['prepared'], y = housing['labels'])`

`forest-reg = RandomForestRegressor()`

~~`forest-reg.fit(x = housing['prepared'], y = housing['label'])`~~

67 Fine Time the model

grid-search. best-param

grid-search. best-estimator

cat-encoder = full-pipeline. named-transformers [cat]

final-model = grid-search. best-estimator

x-test-prepared = full-pipeline. transform
(x=x-text)

final-rmse = np.sqrt (final-mse)

final-rmse

Squared-error = (y-test - final-prediction) ** 2

0.2
1/3 4/10

LAB - 03

Simple Linear Regression and Multiple Linear Regression

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split
```

```
from pandas.core.common import random_state
```

```
from sklearn.linear_model import LinearRegression
```

```
df_sal = pd.read_csv('E:\Salary_Data.csv')
```

```
df_sal.head()
```

```
df_sal.describe()
```

```
plt.title('Salary Distribution Plot')
```

```
sns.distplot(df_sal['Salary'])
```

```
plt.show()
```

```
x = df_sal.iloc[:, :-1]
```

```
y = df_sal.iloc[:, -1]
```

```
X_train, X_test, y_train, y_test = train_test_split  
(x, y, test_size=0.2, random_state=42)
```

regressor = Linear Regression()

regressor.fit(x-train, y-train)

y-pred-test = regressor.predict(x-test)

y-pred-train = regressor.predict(x-train)

print(f'coefficient: {regressor.coef_}')

print(f'Intercept: {regressor.intercept_}')

df_start = pd.read_csv(r'c:\50-startups.csv')

df_start.head()

df_start.describe()

sns.distplot(df_start['Profit'])

plt.show()

regressor = Linear Regression()

regressor.fit(x-train, y-train)

y-pred = regressor.predict(x-test)

np.set_printoptions(precision=2)

result = np.concatenate((y-pred.reshape(len(y-pred), 1),
y-test.reshape(len(y-test), 1),
X))

result

WEEK-4

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
import seaborn as sns

from sklearn.datasets import load_iris

# Load the Iris dataset
iris_data = load_iris()

# Display the dataset description
print(iris_data.DESCR)

# Extract independent features (X) and dependent feature (Y)
X = iris_data.data
y = iris_data.target

# Display the shapes of X and y
print("Shape of X: ", X.shape)
print("Shape of y: ", y.shape)
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split
(X, y, test_size=0.33,
random_state=42)
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
tree_model = DecisionTreeClassifier()
```

```
tree_model.fit(X_train, y_train)
```

```
DecisionTreeClassifier()
```

```
from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt
```

Create a decision tree classifier

```
clf = DecisionTreeClassifier()
```

Train the classifier on Iris Dataset

```
clf.fit(X, y)
```

Plot decision tree

```
plt.figure(figsize=(12, 8))
```

```
plot_tree(clf, filled=True, feature_names=iris_data.
```

```
feature_names, class_names=iris_data.target_names)
```

```
plt.show()
```

Output

petal width(cm) $t = 0.8$

gini = 0.667

samples = 150

values = [50, 50, 50]

class = Setosa

gini = 0.0

samples = 50

values = [50, 0, 0]

class = Setosa

petal width(cm) $t = 1.75$

gini = 0.5

samples = 100

values = [0, 50, 50]

class = versicolor

petal length(cm) $t = 4.95$

gini = 0.168

samples = 54

values = [0, 49, 5]

class = versicolor

petal length(cm) $t = 4.085$

gini = 0.043

samples = 46

values = [0, 1, 45]

class = virginica

from sklearn.model_selection import cross_val_score

scores = cross_val_score(clf, X, y, cv=5)

accuracy = scores.mean()

print("Mean Accuracy:", accuracy)

O/P Mean Accuracy: 0.9666

WEEK - 5

```
import pandas as pd
from matplotlib import pyplot as plt
from sklearn.model_selection import
    train_test_split
from sklearn.linear_model import LogisticRegression
df = pd.read_csv("insurance_data.csv")
print(df.head())
X_train, X_test, y_train, y_test = train_test_split
    (df[['age']], df['bought_insurance'], test_size=0.2)
print("X-test:")
print(X_test)
```

```
model = LogisticRegression()
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
print(y_pred)
```

```
✓ print(model.predict_proba(X_test))
print(model.score(X_test, y_test))
```

Output:

Accuracy = 0.5

import math

def sigmoid(z):

$$\text{return } 1 / (1 + \text{math.exp}(-z))$$

def predi(age):

$$z = 0.042 * \text{age} - 1.53$$

$$y = \text{sigmoid}(z)$$

return y

print(predi(35))

print(predi(43))

Output:

Prediction: array (1, 0, 1, 0, 0, 0, 0, 1, 0)

Score: 0.8888

Linear Reg Slope: 0.584321

Prediction: 0.485

0.5685

Ques
Ans

WEEK-6

```
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
plt.style.use('eggplot')
```

```
df = pd.read_csv('e:/content/diabetes.csv')
```

```
df.head()
```

```
df.shape
```

```
x = df.drop(['outcome'], axis=1).values
```

```
y = df['outcome'].values
```

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split
```

```
(x, y, test_size=0.4,
```

```
random_state=42, stratify=y)
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
neighbors = np.arange(1, 9)
```

```
train_accuracy = np.empty(len(neighbors))
```

```
test_accuracy = np.empty(len(neighbors))
```

for i, k in enumerate(neighbors):

$knn = K\text{NeighborsClassifier}(n_neighbors=k)$

$knn.fit(X_train, y_train)$

$\text{train_accuracy}[i] = knn.score(X_train, y_train)$

$\text{test_accuracy}[i] = knn.score(X_test, y_test)$

`plt.title('K-NN Varying number of neighbors')`

`plt.plot(neighbors, test_accuracy, label='Testing Accuracy')`

`plt.plot(neighbors, train_accuracy, label='Training Accuracy')`

`plt.legend()`

`plt.xlabel('Number of neighbors')`

`plt.ylabel('Accuracy')`

`plt.show()`

`plt.scatter(neighbors, train_accuracy, color='k')`

`plt.scatter(neighbors, test_accuracy, color='g')`

`plt.show()`

$knn = K\text{NeighborsClassifier}(n_neighbors=7)$

~~$knn.fit(X_train, y_train)$~~

~~$knn.score(X_test, y_test)$~~

Accuracy: 0.730519480

(73.05 %)

SVM Model

```
from sklearn import datasets
```

```
cancer = datasets.load_breast_cancer()
```

```
print(cancer.target)
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(cancer.data,  
cancer.target, test_size=0.3,  
random_state=109)
```

```
from sklearn import svm
```

```
clf = svm.SVC(kernel='linear')
```

```
clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
```

```
from sklearn import metrics
```

```
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Output:

✓ Accuracy: 0.9649 (96.49%)

100

W-5.1M
of

WEEK - 7

ANN Program

```
db = np.loadtxt ("..../input/diabetes.csv")
np.random.shuffle(db)
```

y = db[:, 0]

x = np.delete(db, [0], axis=1)

x-train, x-test, y-train, y-test = train-test-split
(x, y, test-size = 0.1)

print(np.shape(x-train), np.shape(x-test))

hidden-layer = np.zeros(72)

weights = np.random.random([len(x[0]), 72]))

Output-layers = np.zeros(2)

hidden-weights = np.random.random([len(x[0]), 72]))

def sum-function (weights, index-col, x):
 result = 0

for i in range (0, len(x)):

result += x[i] * weights[i][index-locked-col]
return result

def activate_layer (layer, weights, x):

for i in range (0, len (layer)):

$$\text{layer}[i] = 1.7159 * \text{np.tanh} (2.0 * \text{sum_function} (\text{weights}, i, x) / 3.0)$$

def back_propagation (hidden_layer, output_layer,
one-hot-encoding, learning_rate, x):

$$\text{output-derivative} = \text{np.zeros}(2)$$

$$\text{output-gradient} = \text{np.zeros}(2)$$

for i in range (0, len (output_layer)):

$$\text{output-derivative}[i] = (1.0 - \text{output_layer}[i]) * \text{output_layer}[i]$$

$$\text{hidden-derivative} = \text{np.zeros}(72)$$

$$\text{hidden-gradient} = \text{np.zeros}(72)$$

for i in range (0, len (output_layer)):

$$\text{hidden-derivative}[i] = (1.0 - \text{hidden-layer}[i]) * (1.0 + \text{hidden-layer}[i])$$

for i in range (0, len (hidden_layer)):

$$\text{Sum} = 0$$

for j in range (0, len (output-gradient)):

sum + = output-gradient [j] * hidden-weight [i]

hidden-gradient [i] = sum * hidden-derivative [i]

recalculate (learning-rate, weights, hidden-gradient, x)

Random Boosting and Ada Boosting

dataset = pd.read_csv ("health-data.csv")

corr-matrix = input-x.corr().abs()

upper = corr-matrix.where (np.sum (np.abs (corr-matrix - shape), k=1) as type (np.bool))

to-drop = [column for column in upper.columns if any (upper[column] > 0.95)]

print (to-drop)

x-train, x-test, y-train, y-test = train-test-split (x, y, test-size=0.1)

n-classes=7

n-estimators=100

RANDOM-SEED=13

names = [`'Random Forest Classifier', 'Ada Boost Classifier'`]

models = [`RandomForestClassifier(n_estimators=n_estimators),`

`AdaBoostClassifier(DecisionTreeClassifier(max_depth=None),`
`n_estimators=n_estimators)`]

for counter model in enumerate(models):

`model.fit(X_train, y_train)`

`y_pred = model.predict(X_test)`

Accuracy Random Forest Classifier : 0.768439 (appx 76.8%)

Accuracy Ada Boost Classifier : 0.7341337 (appx 73.4%)

✓
20/15 fm

WEEK-8

Program 10

K-Means Algorithm

```
iris = pd.read_csv ("---.csv")
```

```
x = iris.iloc[:, [0, 1, 2, 3]].values
```

```
iris_outcome = pd.crosstab (index=iris[["Species"]],  
                           columns="count")
```

iris_outcome

```
sns.FacetGrid (iris, hue="Species").map  
(sns.displot, "petal_length").add_legend()
```

```
sns.FacetGrid (iris, hue="Species").map  
(sns.displot, "petal_width").add_legend()
```

```
sns.FacetGrid (iris, hue="Species").map  
(sns.displot, "Sepal_length").add_legend()  
plt.show()
```

\$

```
sns.set_style ("whitegrid")
```

sns.pairplot (iris, hue = "species", size = 3):
plt.show()

kmeans = kmeans (n_clusters = 3, init = "k-means++",
max_iter = 300, n_init = 10,
random_state = 0)

y_kmeans = kmeans.fit_predict(x)

Program 11

PCA

df = pd.read_csv ("...data.csv")

df_features = df.drop(["diagnosis"], axis = 1)

from sklearn.preprocessing import StandardScaler

Standardized = StandardScaler()

Standardized.fit (df_features)

scaled_data = Standardized.transform (df_features)

from sklearn.decomposition import PCA

pca = PCA (n_components=3)
pca.fit (scaled_data)

x_pca = pca.transform (scaled_data)
scaled_data.shape

x_pca.shape

def diag (x):

if $x = \in \mathbb{M}^n$:
return 1

else: return 0

df_diag = df [c 'diagnosis'].apply (diag)

x_pca [:1]

df_pc = pd. Data Frame (pca.components
columns=df_features.columns)

plt.figure (figsize=(15,8))

sns. heatmap (df_pc, cmap=c 'viridis')

plt.title ('Principal component correlation
with the features?')